

Measuring the MCP Supply Chain: No Elevated Risk on the Dimensions a Registry Exposes, and a Population That Is Not the Population

Prachet Poddar

University of California, Los Angeles

prachetpoddar@gmail.com

Version 1.0.2, 7 September 2026

This work was conducted independently. It was not funded by, supervised by, or carried out under the auspices of any research program, and it should not be attributed to the University of California.

Summary

Model Context Protocol servers are software that runs on developer machines, is launched by unpinned commands, and appears in no dependency manifest. That combination has attracted claims that the MCP ecosystem carries elevated supply-chain risk.

We measured it. Across 250 npm-distributed MCP servers, 2,868 of their transitive dependencies, 400 published packages, and 247 PyPI-distributed servers, we tested eight dimensions against coverage-matched, size-adjusted and age-adjusted controls: license validity, license text distribution, deprecated dependencies, release churn, abandonment, install-script execution, name collision, and license metadata completeness.

No dimension shows elevated risk, and we can now say so without hedging.

Version 1.0.2 enumerates 95.1% of the npm population rather than the 63.8% a score-ordered search frame reaches, which leaves an unidentified remainder of 4.9% instead of 36.2%. MCP servers ship license files more reliably than either control, 23.5% absent against 36.4% and 50.0%, and carry deprecated dependencies no more often. Install-script exposure is a genuine null: the identification bounds contain the control value. These conclusions hold whatever the 405 unreached packages contain.

Enumeration changed one result and rescaled another. Single-version publication, reported as a null in earlier versions, is 45.2% across the population against langchain's 21.2%. Nearly half the published population has published exactly once, and in the deepest stratum it is 95.2%. Release recency stays higher than both controls but falls from 28.5% to 19.5%. The picture is a maintained head on a large dormant tail, and neither is visible from a search frame.

The frame bias is metric-specific. Publishing behaviour moved by up to 20 points under enumeration. Dependency-tree properties moved by less than one: the median tree in the tail holds 96 packages against 97 in the head. Search score predicts how a package is maintained and not what it installs.

The eighth appeared to be a real finding and then dissolved under usage weighting. Python-distributed MCP servers ship with no license information at all at 32.8%

against a PyPI baseline of 18.0%, a risk ratio of 1.8 and an age-adjusted odds ratio of 2.45. Weighted by actual installs, that rate falls to **1.9%**, a sixteen-fold collapse.

The reason is the most useful thing we measured. Downloads in this ecosystem are distributed with a **Gini coefficient of 0.96**, and a single package accounts for **63.6%** of all installs in our sample. The published population and the installed population are almost disjoint. Any study that samples packages uniformly, as every study of this kind including our own first six rounds did, is describing software that essentially nobody runs. Docker's curated catalog, measured the same way, has a Gini of 0.71 and a top item at 11.4%, which suggests the tail is a property of open publication rather than of MCP.

We also report a sampling error that produced a spectacular false positive in our own data, and which we believe affects other work in this area.

1. Why this was measured

An MCP server is typically launched by a line in a client configuration file:

```
{ "command": "npx", "args": ["-y", "@scope/server"] }
```

Three properties follow, and we verified each directly rather than inferring them.

The configuration records no provenance. We retrieved the official `server.json` schema from the MCP registry and searched it. It contains zero occurrences of `license` or `spdx`. The `ServerDetail` object has eleven properties and requires three: `name`, `description`, `version`. The `Package` object carries an optional `version` and an optional `fileSha256`, and neither is required. The client configuration that actually launches the server discards even those.

Nothing is pinned. `npx -y <name>` re-resolves the latest distribution tag on every invocation, into a cache outside the project directory, recorded in no lockfile.

No dependency scanner sees any of it. A manifest-scanning tool pointed at a repository containing such a configuration returns zero components for the entire MCP surface.

To that we add a measured fact. Resolving launch commands and walking the resulting graphs, the median MCP server installs **97 packages** (mean 88.6, p90 165, maximum 589) at depth up to 11. A quarter install exactly one. Two thirds install more than 50.

The blind spot is real and its contents are numerous. This study asks a narrower question: **is the software in the blind spot unusually bad?**

2. Method

2.1 Sampling frames

npm. Version 1.0.2 enumerates 95.1% of this population. The three paragraphs below describe the original 63.8% frame, which every result in versions 1.0.0 and 1.0.1 rests on; Section 2.5 describes how the rest was reached and what it changed.

The registry search API reports 8,227 packages carrying the keyword `mcp-server`. Pagination stops near 5,000 results, giving a frame of 5,250 unique packages, or 64% coverage. Because search orders by score, this frame is a top-slice and our npm estimates are optimistic. We tested for a hygiene gradient across score quartiles and found none (point-biserial $r = -0.006$ for unevaluable licenses, $r = -0.059$ for missing license files, against a significance threshold of 0.098 at $n = 400$).

PyPI. The complete PyPI simple index is a single request returning **886,346 projects**. Filtering yields 18,963 packages carrying an mcp token and 4,062 that also contain "server". Treatment and control arms are simple random samples from this complete enumeration, with no score ordering.

Controls. After an initial control failed (Section 6), we used coverage-matched npm frames:

Frame	Population	Our frame	Coverage
keywords:mcp-server	8,227	5,250	64%
keywords:eslint-plugin	7,075	5,250	74%
keywords:langchain	1,973	1,973	100%

2.5 Enumerating the population

Free-text term expansion lifted coverage from 63.8% to 79.5% and then saturated: the last ten of a hundred query terms added 26 packages. What broke the ceiling was partitioning on maintainer. `keywords:mcp-server maintainer:cyanheads` returns 121 against the unrestricted 8,238, so the qualifier is applied. Every package has a maintainer and no maintainer's slice approaches the 5,000 pagination cap, so complete maintainer coverage implies complete package coverage.

`scope:` is not a partition. `keywords:mcp-server scope:pipeworx` returns 8,238, identical to no filter at all. npm accepts the qualifier and ignores it. A frame built on it would look partitioned and would not be, which is the same failure as the Docker ordering parameter in Section 6.

The crawl has one blind spot of its own. Maintainers are discovered from search results and search is score-ordered, so a maintainer whose every package sits below the cutoff is never seen. Reading the full maintainers array from each packument, rather than the single publisher search returns, surfaced 1,415 maintainers search never showed. That is what took coverage from 82% to 95.1%. One of them held 1,286 packages.

Final coverage is 7,833 of 8,238, in four strata:

Stratum	Reached by	Size	Share
A	score-ordered search	5,250	63.7%
B	free-text term expansion	1,290	15.7%
C	maintainer partitioning	1,293	15.7%
D	nothing	405	4.9%

Strata B and C were sampled at $n = 250$ each and measured with the same code as stratum A. Population estimates are the size-weighted combination of the three observed strata. Bounds set stratum D to all-clean and all-bad, so a comparison is identified when the bound does not reach the control value.

2.2 Dependency resolution

We resolved npm semantic-version ranges using an implementation of npm range semantics rather than resolving every edge to latest. Distribution tags are consulted first, then maximum-satisfying resolution over published versions. Optional dependencies are included; peer dependencies excluded. Non-registry specifiers are classified rather than dropped.

Resolution correctness matters. Re-walking 60 servers with correct range resolution against naive latest resolution changed the node count on **72% of them**, with a net increase of 15.5% in total package instances.

2.3 Usage weighting

Download counts come from the last-month PyPI statistics endpoint, retrieved for 221 of the 247 sampled packages; 26 returned no record.

2.4 Statistics

Medians of even-sized samples are reported as the upper of the two central order statistics, so a sample whose exact median is 96.5 is reported as 97. Proportions carry 95% normal-approximation intervals; comparisons use a two-proportion z test. Where an outcome scales mechanically with dependency-tree size or package age, we report Mantel-Haenszel odds ratios stratified on that confounder and restrict both arms to a common range.

3. Result: extreme usage concentration

This is reported first because it governs the interpretation of everything else.

Across 221 PyPI-distributed MCP servers with 203,446 combined monthly downloads:

Measure	Value
Gini coefficient	0.9611
Share taken by the single top package	63.6%
Share taken by the top 5	88.7%
Share taken by the top 10	93.0%
Share taken by the top 25	96.5%
Median downloads, top decile	687/month
Median downloads, bottom decile	8/month

One package accounts for roughly 129,000 of 203,000 monthly installs. The bottom half of the sample is measured in single-digit downloads per month.

An ecosystem of several thousand published packages is, in practice, a few dozen packages. Uniform sampling over the published population, which is what every measurement in Sections 4 and 5 does, describes software that is published but not used.

3.1 The curated channel has a different shape

Docker distributes MCP servers through a hand-maintained catalog rather than an open registry. We took all 328 entries from that catalog and fetched pull counts for each repository individually. 182 have a published image; the remaining 146 are catalogued but were never published.

Measure	PyPI (open registry)	Docker mcp/ (curated catalog)
Packages with usage data	221	182
Gini coefficient	0.9611	0.7094
Top 1 share	63.6%	11.4%
Top 5 share	88.7%	38.6%
Top 10 share	93.0%	53.3%
Top 25 share	96.5%	69.9%
Items with zero usage	bottom decile at 8/month	0

The two figures are not comparable in absolute terms. Docker reports lifetime pulls, PyPI reports one month, and both include CI and mirror traffic. Only the shape is comparable, and the shape differs. The curated channel is skewed but not degenerate: its top item takes 11.4% rather than 63.6%, its median image has 33,892 pulls, and no image in it has zero.

This is what curation does to a distribution. It is not evidence that Docker's servers are better; the catalog was assembled by selecting servers people already wanted. The point is that the long tail of unused packages, which is what dissolved the Section 5

finding, is largely a property of open publication rather than of MCP. A channel that publishes 328 things by hand does not have one.

It also gives the allowlist argument in Section 9 a concrete size. Docker has already built such a list, and it is 328 names long, of which 25 carry 70% of the pulls.

4. Result: no dimension of elevated risk

All seven npm dimensions are measured against the two coverage-matched frames described in Section 2.1. Version 1.0.0 of this paper measured rows 1 to 4 against a keywords:cli plus keywords:server frame instead, which Section 6 shows to be invalid; those four rows have been re-measured and two of them changed direction. Section 11 records what changed.

#	Dimension	MCP, stratum A only	MCP, coverage-corrected	Identification bounds	eslint-plugin	langchain
1	Ships no license file	22.2%	23.5%	[22.3, 27.2]	36.4%	50.0%
2	Declares a license, ships no text	21.2%	22.0%	[20.9, 25.9]	36.0%	49.2%
4	Released in last 30 days	28.5%	19.5%	[18.5, 23.5]	11.6%	16.4%
5	Published exactly one version	25.2%	45.2%	[43.0, 47.9]	26.4%	21.2%

Rows 3, 6 and 7 are tree-shaped or name-shaped rather than metadata-shaped and are treated separately below. Corrected the same way, they move very little:

#	Dimension	Stratum A	Tail B+C	Corrected	Bounds	eslint	langchain
3	Deprecated package in tree	17.2%	20.1%	18.1%	[17.3, 22.2]	20.4%	25.6%
6	Install hook in tree	15.6%	13.3%	14.8%	[14.1, 19.0]	not measured	14.2%

The bias is metric-specific, and that is the most useful thing enumeration taught us. Single-version publication moved by 20 points and release recency by

1. Deprecated dependencies moved by 0.9 and install hooks by 0.8. The reason is visible in the trees: the median dependency tree in the enumerated tail has 96 packages against 97 in the score-ordered head. A package that published once and stopped installs the same amount of code as one that publishes weekly. Search score predicts how a package is maintained; it does not predict what it pulls in.

Against langchain, which is enumerated completely and needs no correction, all four bounds clear the control value. Those four comparisons are identified: they hold whatever the 405 unreached packages contain. Against eslint-plugin, which is itself a 74% frame and therefore biased in the same direction, only rows 1, 2 and 4 are identified; row 5 is not, because a corrected eslint-plugin figure could reach 45.5%.

Two rows did not survive enumeration.

Row 5 was reported as a null in versions 1.0.0 and 1.0.1. It is not. Single-version publication is 45.2% across the enumerated population against langchain's 21.2%, and in stratum C alone it is 95.2%. Nearly half the published MCP population has published exactly once. The null was an artifact of measuring the maintained top-slice of one ecosystem against a fully enumerated other.

Row 4 moved the other way. The naive 28.5% becomes 19.5% once the tail is included, because the tail is almost entirely dormant: 0.4% of stratum B and 2.0% of stratum C published in the last 30 days. The direction survives, the magnitude does not.

No row shows MCP servers worse than a coverage-matched control on licensing or deprecation. Two rows show them differing on maintenance, and the difference is in both directions depending on which maintenance measure is used: they publish more recently and they are far more likely to have published only once. Both are consistent with an ecosystem where a maintained head sits on a large dormant tail.

Licensing (rows 1 and 2). Version 1.0.1 reported these as bounds because the MCP frame was 64% and score-ordered while langchain was 100%. At 95.1% enumeration the hedge is no longer needed. The corrected rates are 23.5% and 22.0% against 36.4% and 50.0%, and the identification bounds reach only 27.2% and 25.9%. MCP servers ship license files more reliably than either control, and that conclusion does not depend on what the 405 unreached packages contain.

Deprecated dependencies (row 3). At 18.1% corrected, with bounds reaching 22.2%, MCP is identified as lower than langchain's 25.6% and not identified against eslint-plugin's 20.4%. No control satisfies both requirements here. The coverage-matched frames have median tree sizes of 8 and 10 packages against MCP's 97, and a one-package tree cannot contain a deprecated dependency, so they understate the control. The shape-matched cli/server frame has a median of 13 and a mean of 79, closer on shape but invalid on coverage; against it the tree-size-stratified odds ratio is 0.50. Both comparisons point the same way and neither is clean. What can be said is that MCP servers carry deprecated dependencies no more often than any control we have, despite installing more packages: ten times as many by median, though the control distributions are heavily right-skewed and by mean the langchain gap is only 1.6-fold.

Maintenance (rows 4 and 5). These two moved the most and they move in opposite directions, which is the useful part.

Version 1.0.0 reported MCP publishing less often than the control, $z = -4.15$, against the invalid frame. Version 1.0.1 reversed that against coverage-matched frames. Version 1.0.2 keeps the direction but cuts the magnitude: 19.5% rather than 28.5%, because the enumerated tail is dormant, publishing in the last 30 days at 0.4% and 2.0% in strata B and C.

Single-version publication went the other way and is the largest correction in the paper. Stratum A says 25.2%. Stratum C says 95.2%. The population says 45.2%. A frame ordered by search score is, for this metric, close to a filter on the outcome being measured.

Together they describe one shape rather than two findings: a maintained head that publishes often, and a large dormant tail that published once and stopped. Both are visible only if you enumerate.

Name collision (row 7), re-derived. The figures given in versions 1.0.0 and 1.0.1, 14.1% against 9.1% and 20.7%, are withdrawn. The code that produced them was never released, so they cannot be reproduced by us or by anyone else, and a reimplementations from the paper's own description of the metric does not reproduce them.

The metric is therefore redefined explicitly. Strip any scope, lowercase, remove every character that is not a letter or a digit, then remove the ecosystem's own tokens; two packages collide if the remaining stem is identical. Collision rate is strongly coverage-dependent, since a name can only be seen to collide with a name you have, so all three arms are measured at the lowest coverage available to any of them, 74.2% of each population, averaged over 25 random draws.

Arm	Coverage available	At matched 74.2% coverage
MCP	95.1%	16.8% $\pm$ 0.6
eslint-plugin	74.2%	9.4%
langchain	100%	19.6% $\pm$ 1.6

MCP sits between the two, which is what versions 1.0.0 and 1.0.1 concluded from figures that cannot be reproduced. The conclusion is unchanged and now rests on a stated definition and a matched comparison. At its own full 95.1% coverage the MCP rate is 19.2%, which is the figure to use when asking how much name reuse exists rather than how MCP compares.

Install scripts (row 6). The dimension we most expected to differ, because `npx -y` executes `preinstall`, `install` and `postinstall` for the root and every dependency. The absolute exposure is real: **15.6% of MCP server installations execute third-party code at install time, rising to 57.7% once the tree exceeds 120 packages**, usually through native-compilation packages such as `protobufjs`, `sharp`, `better-sqlite3` and `onnxruntime-node`. But `langchain` packages show 14.2%, the size-adjusted odds ratio is 0.79, and the corrected MCP rate of 14.8% has bounds of [14.1, 19.0] that straddle it. This is the one row in the paper that is a null in the strict sense: the identification bounds contain the control value, so no enumeration effort available to us would separate them.

4.1 Why the licensing dimensions do not separate

License metadata on npm is saturated. Within 15 functional clusters the clean-SPDX rate runs 85% to 100% in both head and tail. Pooled, the 20 highest-ranked packages per cluster are 93.7% clean against a random tail at 98.7% ($z = -3.19$), so popularity selects weakly *against* license quality.

A variable sitting at 95% regardless of slicing cannot produce a differential result. npm init -y writes "license": "ISC" into every scaffolded package with no author decision, which sets the ceiling. It writes no LICENSE file, which is the direct cause of the rate at which npm packages of all kinds declare a license they do not ship. That rate is not one number: it is 21.2% for MCP servers, 36.0% for eslint-plugin packages and 49.2% for langchain packages. Version 1.0.0 gave it as roughly one in four, which was the withdrawn control's 24.8%.

5. Result: a finding that does not survive usage weighting

Python-distributed MCP servers are considerably more likely than the PyPI baseline to carry no license information in any of license, license_expression, or a license classifier.

	MCP servers	Control	
n	247	245	
No license information at all	32.8% ±5.9 (81/247)	18.0% ±4.8 (44/245)	$z = +3.78$
Median age since first upload	312 days	1,123 days	
Age-adjusted odds ratio			2.45
Both arms restricted to ≤600 days	32.7% ±5.9	14.5% ±7.9 (n=76)	$z = +3.07$

It survives age matching, the confound that destroyed two of our npm results. Our general-PyPI control rate of 18.0% agrees within its interval with the 15.41% reported for PyPI by a 2024 study of 33.7 million versioned packages, so the control measures what independent work says it should.

Then it collapses under usage weighting.

	Rate
Package-weighted	31.7% (95% CI 25.9 to 38.1), 70/221
Download-weighted	1.9% , 3,941 of 203,446 installs
Collapse factor	16.4x

The decile breakdown shows why. Missing licenses are concentrated almost entirely in packages nobody installs:

Decile	Median downloads/month	No license info
1 (most downloaded)	687	18.2%
2	175	4.5%
3	63	4.5%
4 to 6	38 to 18	22.7%
7	14	36.4%
8	13	45.5%
9 and 10	11 and 8	68.2%

Top three deciles average 9.1%; bottom four average 54.6%.

A prediction we got wrong, and it matters. Before obtaining download data we used release count and recency as adoption proxies. Packages with four or more releases and activity in the past year showed 31.1% missing, essentially identical to the overall rate, and we predicted the download-weighted figure would land between 20% and 35%. It is 1.9%. **Maintenance activity is not a proxy for adoption.** A package can publish ten releases and be installed eight times a month.

The finding is therefore true and unimportant: real for the published population, close to absent for the installed one.

6. A sampling error that produces false positives

Our first control set produced a spectacular and entirely false result.

Comparing MCP servers against npm packages keyworded `cli` or `server`, single-version publication appeared at 25.2% for MCP against 3.2% for the control, $z = +7.32$. Age was not the explanation: restricting both arms to packages under 561 days old preserved the effect at an age-adjusted odds ratio of 16.86.

The explanation was coverage. `keywords:cli` returns 104,818 packages and we had sampled the top 2,500 by search score, the top **2.4%**. Our MCP frame was 64% of a much smaller population. We had compared most of a small ecosystem against the elite of a large one.

Against coverage-matched frames the effect vanishes: `eslint-plugin` 26.4%, `langchain` 21.2%, MCP 25.2%, age-adjusted odds ratios 1.84 and 1.09. **Roughly a quarter of npm packages publish exactly once. That is what npm looks like when you do not select on popularity.**

Undetected, this error invalidated four earlier results before we caught it. Any study drawing an npm frame from the search API inherits it. We recommend matching

coverage fractions explicitly across arms, or drawing frames from a source that enumerates completely, as the PyPI simple index does.

We then made the same mistake again, in the same paper. Version 1.0.0 rebuilt rows 5 to 7 on coverage-matched frames and left rows 1 to 4 on the invalid one, while Section 2.1 described the whole table as coverage-matched. Re-measuring those four rows changed two of them, and one changed sign: release churn was reported as MCP publishing significantly less often, and against matched frames MCP publishes significantly more often. The lesson we would draw from our own case is that identifying a sampling error is not the same as removing it, and that the rows you do not re-run are the ones that keep it.

7. Incidental result: functional redundancy

Classifying all 5,250 frame packages by declared function:

Cluster	Servers	Cluster	Servers
memory/knowledge	302	browser automation	124
github/git	297	maps/weather	97
database/sql	293	slack/chat	90
finance/crypto	270	notion/docs	87
aws/cloud	163	web search	77
web fetch/scrape	135	filesystem	76
email	130	time/date	44
		jira/issues	42

Heavy redundancy: 297 packages address GitHub, 293 address databases, 76 address the filesystem. **Thirteen packages are effectively "the MySQL MCP server"**, one unscoped as `mcp-server-mysql`, which reads as canonical. Across collision groups, 126 of 274 contain an unscoped name and 62% show a rank gap greater than 1,500 places between their most and least prominent member. That is the shape typosquatting exploits, though the collision rate itself sits between langchain's 20.7% and eslint-plugin's 9.1%.

12 of 15 clusters have a cleanly licensed option in every one of their five most prominent members, and the remaining three have four, three and four of five. All 15 have at least one clean option.

8. Limitations

The npm frame covers 95.1% of its population. The excluded portion does differ categorically on maintenance measures, which is why versions 1.0.0 and 1.0.1 were wrong about two of them, and does not differ on dependency-tree measures. The 405

packages no method reached are unobserved by us and by every published study of this ecosystem we are aware of. This matters more than it did in version 1.0.0: our control frames cover 74% and 100%, so the residual mismatch runs in MCP's favour on every dimension where prominence predicts hygiene. Every MCP figure in Section 4 is now a coverage-corrected estimate with identification bounds attached, computed over three observed strata with the unreached 4.9% set to both extremes. The control arms are not corrected the same way: langchain is enumerated completely and needs none, but eslint-plugin is a 74% frame and is biased in the same direction MCP's was, so every eslint-plugin comparison should be read as provisional and only the langchain comparisons are identified.

Two figures in this paper cannot be recomputed from the release. The score-quartile hygiene-gradient test in Section 2.1, $r = -0.006$ and $r = -0.059$, needs a per-package search rank that `npm_stage1_metadata` does not carry; that test is the only defence of our 64% score-ordered frame, which version 1.0.1 leans on more heavily than 1.0.0 did. And the age-adjusted odds ratio of 1.84 in Section 6 is sensitive to the choice of age strata: standard bins give values near 1.0, and 1.84 is reachable only under the 561-day common-support restriction. The direction of that sensitivity is conservative, since a value nearer 1.0 strengthens rather than weakens the claim that the effect vanishes.

The released files disagree on one denominator. `mcp-confound-tests.json` records the PyPI control as 44 of 245, giving 18.0%; `mcp-controls-data.json` records 44 of 246, giving 17.9%. Section 5 uses 245. The difference moves z from +3.78 to +3.80 and changes nothing.

No available frame is matched on both coverage and dependency-tree shape. The coverage-matched frames have median trees of 8 and 10 packages against MCP's 97; the shape-matched frame is a 2.4% top-slice. Row 3 is reported against both and is clean against neither.

Download data was obtained for PyPI only. The npm arm remains package-weighted, though Section 4.1 shows license quality does not vary with popularity there, so weighting would move those results little.

The download sample is 221 packages from a 4,062-package frame at one point in time, and last-month counts include mirror and CI traffic, which inflates absolute figures without obviously biasing the comparison.

Peer dependencies are excluded, so trees are undercounts. Docker-distributed servers are measured for usage concentration only (Section 3.1); their images were not inspected for license or maintenance metadata, so they contribute to no result in Sections 4 or 5.

We measured license and maintenance metadata. We did not measure malware, typosquatting outcomes, prompt injection, tool poisoning, or runtime capability scope, which is where most published MCP security work sits. **Our results say nothing about those dimensions.**

9. What this means

For anyone assessing MCP supply-chain risk on the dimensions a package registry exposes: this ecosystem is not worse than its neighbours, and on licensing it is measurably better than both of our controls, with identification bounds that do not reach either. Claims that it is unusually poorly licensed, unusually abandoned, or unusually prone to install-time code execution are not supported once coverage-matched controls are applied. The one dimension where a real difference exists affects 1.9% of installs. MCP servers do publish more often than comparable packages, which is a sign of activity rather than of risk.

The larger methodological point is that **the published population is the wrong unit of analysis**. With a Gini of 0.96 and a single package taking 63.6% of installs, per-package rates describe a long tail that is measured but not run. This applies to our own npm results as much as to the finding it dissolved, and it applies to any registry-scale study that samples uniformly.

There is a practical consequence. If a few dozen servers account for nearly all real usage, the governance problem is far smaller than package counts suggest, and a curated allowlist covering the top few dozen would address most deployments.

The structural facts stand independently of all of this. Agent tool configurations record no provenance, nothing is pinned, a median of 97 packages executes per server with 15.6% running install-time code, and no existing scanner sees any of it. Whether that absence of visibility is a problem anyone needs solved is not a question measurement can answer.

10. Data and code

Availability. Every measurement, sampling frame, control arm and analysis script is released at:

<https://github.com/prachetpoddar/mcp-supply-chain-study>
DOI 10.5281/zenodo.22641813 archives **version 1.0.0**, which contains the errors listed in Section 11. Version 1.0.1 is archived under its own DOI, minted from the v1.0.1 tag; the Zenodo concept DOI always resolves to the newest version.

The repository contains sixteen analysis scripts plus the resolver library, and twenty-four data files covering all four rounds of measurement, including the two runs that were later found invalid and the corrections applied to them. Nothing has been removed to make the result look cleaner.

Four files were added in 1.0.1 to close reproducibility gaps a reader found in 1.0.0: `pypi_weighted_result.json`, which carries the per-package download counts behind every figure in Section 3; `walk250.jsonl`, which carries package-level attribution for the 250 MCP dependency trees and is what makes row 3 recomputable; `rerun_state_raw.json`, the raw tarball and walk records for the 500 re-measured control packages; and `rows_1_4_rematched.json`, the summary built from them, alongside `rerun_rows_1_4.py` which produced both.

Licensing. Code is released under Apache-2.0. Data is released under CC-BY-4.0. The datasets consist of factual metadata retrieved from public registry APIs, together with derived statistics; no third-party package source is redistributed.

Provenance. The resolver uses only the public npm and PyPI registry APIs and the public Docker Hub repository API. Package trees were resolved with npm semver range semantics rather than by taking latest, which is why the round-one dataset is superseded. License detection uses the SPDX License List, 739 identifiers at time of measurement. Docker pull counts were taken one repository at a time from the 328-entry docker/mcp-registry catalog, because the namespace listing endpoint silently ignores its ordering parameter (Section 6).

Reproduction notes. Three things will otherwise reproduce our mistakes rather than our results.

1. Do not build a sampling frame from the npm search API without matching coverage fractions across arms. Section 6 gives the false positive this produces, $z = 7.32$ for an effect that does not exist.
2. Verify that an ordering parameter took effect before building anything on the result. Docker Hub accepts `ordering=-pull_count` and ignores it.
3. Do not skip a request that returns HTTP 429. Rate limiting is correlated with alphabetical position, which is correlated with nothing, until it is: our first Docker run lost the entire s-to-z tail, including six of the ten largest repositories.

Weight by usage before drawing conclusions about a population. That is the one instruction in this paper we would keep if we could keep only one.

11. Changes since version 1.0.0

Version 1.0.0 was published on 7 September 2026 and archived under DOI 10.5281/zenodo.22641813. An independent check of the paper's numbers against the released data, run the same day, found the following. Every change is listed, including the ones that weaken the paper.

Rows 1 to 4 were measured against an invalid control and have been re-measured. Section 2.1 described the table as coverage-matched. Only rows 5 to 7 were. Rows 1 to 4 used the keywords:cli plus keywords:server frame that Section 6 shows to be a 2.4% top-slice. Re-measured against the coverage-matched frames:

Row	1.0.0	1.0.1
1 Ships no license file	null, $z = -1.09$	MCP lower, $z = -3.92$ and -7.32
2 Declares, ships no text	null, $z = -1.05$	MCP lower, $z = -4.12$ and -7.43
3 Deprecated in tree	null, $z = -1.57$	null and MCP lower, $z = -0.92$ and -2.29
4 Released in last 30 days	MCP lower , $z = -4.15$	MCP higher , $z = +5.06$ and $+3.52$

Row 4 changed sign. The 1.0.0 claim that MCP servers publish less often than comparable npm applications was an artifact of the invalid frame and is withdrawn.

The row 4 control proportion in 1.0.0 was also wrong. It was reported as 111 of 250. The per-package records in `mcp-confound-tests.json` give 104 of 250 and $z = -3.44$ rather than -4.15 . The MCP numerator, 114 of 400, reproduces exactly. Note that the summary block in `mcp-controls-data.json` still asserts 111; it is shipped uncorrected so that the discrepancy is visible rather than tidied away.

Row 7 was printed with the z column blank. The underlying measurement records $z = +8.01$ against `eslint-plugin` and $z = -6.83$ against `langchain`. Both are now shown. The row is still reported as a null because MCP sits between two comparable ecosystems, but readers should see the magnitudes.

The headline was overstated. 1.0.0 was titled "Eight Null Results" and claimed seven nulls. Under corrected controls the count is two clean nulls, one row that is null against one control and favourable against the other, one row that sits between two ecosystems, two rows where MCP is lower, and one where it is higher. The conclusion that no dimension shows elevated risk is unchanged and is now better supported. The word "null" is no longer accurate and has been removed from the title.

A limitation was understated. Our MCP frame covers 64% against controls at 74% and 100%, so the residual mismatch runs in MCP's favour. Rows 1, 2 and 4 are now reported as bounds. Separately, no frame is matched on both coverage and tree shape, which affects row 3.

Section 3 was not reproducible. The per-package download counts behind the Gini of 0.9611, the 63.6% top-1 share and the 1.9% weighted rate were not in the release, while Section 10 claimed everything was. `pypi_weighted_result.json` has been added. So has `walk250.jsonl`, without which row 3 could not be recomputed.

Row 7 was called a null and is not one. It is two significant results in opposite directions. Version 1.0.1 as first drafted repeated 1.0.0's framing; this was caught by a second verification pass and the row is now reported as a bound. The title gained the qualifier "on the dimensions a registry exposes" in the same pass, because "no elevated risk" unqualified asserts more than the table supports.

Prose downstream of the corrected rows was not corrected in the first pass. Section 4.1 still quoted a one-in-four rate for npm packages declaring a license they do not ship, which was the withdrawn control's 24.8%. The release README still carried the odds-ratio-as-risk-ratio phrasing and the 8,229 figure that this changelog says were fixed. Both are now corrected. This is the third time in one paper that a correction was applied in one place and not the others.

Four smaller corrections. The npm frame population is 8,227, not 8,229, stated twice. Maximum dependency-tree depth is 11, not 10; the figure of 10 came from the superseded round-one dataset. The functional-cluster count is 12 of 15 with a clean option in all five most prominent members, not 13, and the remainder is three clusters at four, three and four of five, not two at three of five. The summary described an odds ratio of 2.45 as "2.4 times more likely", which is a risk-ratio phrasing; the risk ratio is 1.8.

None of these changes affect Section 3, Section 5, Section 7, or the structural facts in Section 1. The PyPI licensing result, the usage-concentration result, and the install-hook figures all reproduce from raw records unchanged.

12. Changes since version 1.0.1

Version 1.0.1 reported four rows as identification bounds because the npm frame covered 63.8% of its population and was drawn in search-score order. Version 1.0.2 enumerates 95.1% instead, which replaces the hedging with measurement.

Enumeration. Maintainer partitioning plus reading the full maintainers array from every packument took coverage from 63.8% to 95.1%, 7,833 of 8,238 packages. Section 2.5 gives the method, including the two things that did not work: free-text term expansion saturates near 79.5%, and the scope: qualifier is accepted by npm and silently ignored.

Row 5 was not a null. Single-version publication is 45.2% across the population, not the 25.2% the score-ordered frame showed, against langchain's 21.2%. The identification bound is 43.0%, so the comparison holds regardless of the unreached remainder. Versions 1.0.0 and 1.0.1 both called this a null.

Row 4 keeps its direction and loses half its size. Release recency falls from 28.5% to 19.5% against controls at 11.6% and 16.4%.

Rows 1 and 2 stop being bounds and become findings. Corrected rates of 23.5% and 22.0% against 36.4% and 50.0%, with bounds reaching only 27.2% and 25.9%.

Rows 3 and 6 barely move, 17.2% to 18.1% and 15.6% to 14.8%. The median dependency tree in the enumerated tail holds 96 packages against 97 in the head, so search score predicts maintenance behaviour and not installed footprint. Row 6 is the paper's one strict null: its bounds contain the control value.

Row 7 is withdrawn and re-derived. The 14.1% / 9.1% / 20.7% figures in versions 1.0.0 and 1.0.1 came from code that was never released and cannot be reproduced, including by us. Section 4 now states the metric explicitly and measures all three arms at matched coverage: 16.8%, 9.4% and 19.6%. The conclusion, that MCP sits between two comparable ecosystems, is unchanged.

What did not change. Section 3, Section 5 and Section 7 are untouched. The PyPI arm was drawn from the simple index, which enumerates completely, so it never had this problem. The usage-concentration result and the structural facts in Section 1 stand as published.

Standing caveat. The control arms have not been enumerated to 95%. eslint-plugin remains a 74% frame carrying the same bias MCP's frame had, so those comparisons are provisional. Only the langchain comparisons are identified.